

Ragul Narayanan Magesh

Boston, MA | magesh.ra@northeastern.edu | [LinkedIn](#) | [Portfolio](#) | [GitHub](#) | +1 (617) 792-7068

Education

Northeastern University

May 2026

Master of Science in Data Analytics Engineering (Data Science & Machine Learning Focus) | **GPA 3.89**

Boston, MA

Relevant Coursework: Applied Generative AI, Deep Learning, Data Mining, ML Operations, Database Management, NLP

Skills and Certifications

Programming: Python (Pandas, NumPy, Scikit-learn, PyTorch, TensorFlow), SQL (Window Functions, CTEs), PySpark, Linux

GenAI & LLMs: LLMs, RAG, LangChain, Prompt Engineering, Fine-Tuning, Agentic AI, OpenAI API, HuggingFace, MCP

MLOps & Cloud: Claude Code, Docker, Apache Airflow, MLflow, DVC, CI/CD, Git, AWS (S3, SageMaker), GCP (VertexAI)

Data: Snowflake, PostgreSQL, MongoDB, Firestore, FastAPI, Streamlit, Tableau, PowerBI

ML & Deep Learning: Supervised, Unsupervised, Deep Learning, CNNs, XGBoost, Random Forest, Hyperparameter Tuning

Certifications: Claude Code 101, Claude Code in Action, Model Context Protocol Advanced Topics (Anthropic)

Experience

KCC Capital Partners

Jan 2026 - Apr 2026

Data Scientist Intern - AI Engineering

Remote, United States

- Built RAG retrieval pipelines for the Donna AI platform using LangChain, vector databases, and OpenAI on GCP, supporting deal sourcing and document processing workflows across 30 specialized LLM agents
- Developed the legal drafting agent's retrieval system by tuning prompts and evaluating outputs across OpenAI and Gemini models, improving accuracy of NDA and term-sheet generation
- Collaborated with business stakeholders to identify AI use cases, translating investment banking requirements into functional prototypes and documenting pipeline configurations for team-wide adoption

360DigiTMG

Aug 2023 - Jul 2024

Data Scientist

Chennai, India

- Analyzed IQVIA pharmaceutical sales data to evaluate product performance and benchmark against competitors across multiple therapeutic categories
- Transformed and modeled datasets in Snowflake using advanced SQL, improving reporting efficiency by 40% and enabling seamless integration with Tableau dashboards visualizing sales trends, prescription activity and regional performance
- Performed predictive analytics and ML-based HCP profiling to forecast sales, identify high-value prescribers, and optimize marketing and sales targeting, resulting in a 25% increase in targeting precision

Wipro

Apr 2022 - Jul 2023

Data Engineer

Bengaluru, India

- Migrated 150+ Teradata objects containing retail, customer, and supply-chain data into Snowflake, ensuring metadata alignment, validating data consistency and optimizing structure for downstream analytics
- Built Python-based semantic monitoring tools to validate data quality, enforce business rules, and trigger real-time alerts, ensuring system reliability across production pipelines
- Optimized ETL pipelines using PySpark and Airflow, cutting processing time by 30% and automating incremental data loads, enabling timely delivery of analytics-ready datasets

Projects

MarketMind - AI-Powered Daily Stock Briefing Platform | [Live Demo](#)

Apr 2026

- Built a 5-agent LLM pipeline orchestrated via LangChain that reads overnight financial news, SEC filings, and market data to generate a personalized daily brief for investors using Claude Code as the primary development tool
- Implemented multi-modal analysis covering news sentiment scoring, SEC filing RAG, real-time financial metrics, and macro event correlation, synthesized into recommendation signals with confidence scores and narrated audio briefs
- Deployed on GCP Cloud Run with a daily scheduler to scrape news articles and store embeddings in Qdrant Cloud for semantic search, Firestore for user data and portfolio persistence, and OpenAI TTS for audio generation

EchoAI - End-to-End MLOps Pipeline for Automated Review Response

Dec 2025

- Built automated data pipelines using Cloud Composer and BigQuery to scrape, merge, preprocess, and feature engineer customer review data from multiple sources, feeding 100K+ records into downstream model training and evaluation
- Automated model training using Airflow and MLflow to train and compare sentiment models, select the best performer based on accuracy, precision, recall, and F1-score, and continuously monitor for bias and data drift to trigger retraining
- Deployed a FLAN-T5 powered response generation application on GCP Cloud Run with a Streamlit UI supporting single review and bulk CSV uploads with confidence scores, containerized with Docker and CI/CD (GitHub Actions)